

# Reinforcement Learning Project

Raphael Monges

**Topic:** Stock Trading - We choose a commodity with volatility and seasonality: Gaz.

## Motivation :

As being a student who specialized in financial markets and who is doing his internship in it, I am motivated to develop tools with reinforcement learning to improve algorithmic trading strategies through reinforcement learning and if time available, reinforcement learning with human feedbacks.

## Description of the environment:

### State space:

The state space in a reinforcement learning framework for stock trading can be mathematically defined as a vector of features  $s_t \in S$  representing the market and internal state at time  $t$ .

This includes:

1. Price vectors  $P(t,i)$  that represents close price at time  $t$  of stock  $i$ .
2. Technical indicators  $I(t,i)$  that represents an indicator  $I$  at time  $t$  for stock  $i$ .
3. Other features  $A_t$ , such as CPI, interest rates from central bank, positions, at time  $t$

Moreover, the state at time  $t$  can be represented as :  $s_t = [P_t, I_t, A_t]$

### Action space:

All the possible actions  $a_t \in A$  that the agent can take at each timestep are =

1. Buy  $x$  shares of stock  $i$  such as  $a(t,i) = x > 0$
2. Sell  $x$  shares of stock  $i$  such as  $a(t,i) = x < 0$
3. Hold:  $s(t,i) = 0$

(We do not take into account hedging here with options for example)

### Reward space:

The reward function  $R(s_t, a(t,i))$  provides feedback to the agent based on the outcome of its actions guiding the learning process.

$R(s_t, a(t,i)) = \text{VariationPV}(t) - \lambda$  where  $PV(t)$  is the change in portfolio value resulting from action  $a(t)$  at state  $s(t)$  and  $\lambda$  is a risk aversion coefficient representing volatility or drawdown.

### Description of the implemented agent:

As RL algorithm, at first we'll develop a Q-learning and we are looking forward to use in a continuous space the Deep Q-Networks. We want to learn a policy  $\pi^*$  that max expected returns. Then we would have:

Q-learning :  $Q^*(s,a) = E[R(s,a) + \gamma \max_{a'}(Q^*(s', a'))]$

DQN: an approximation of  $Q^*(s,a)$  using a neural network

If time available we would be able to make the agent learn from Human Feedback, it implies modifying the reward function adding a term:

$$R'(s_t, a_t) = R(s_t, a_t) + n * H(s_t, a_t)$$

### Conclusion:

Our goal is to optimize the policy  $\pi^*$  such as it maximizes the expected cumulative reward, taking into account market returns and market risk. The challenge will be to include a stochastic model such as  $[P(s_{t+1}|s_t, a_t)]$  that generalizes well to unseen market conditions.